

# Introduction to Supervised Machine Learning


From Linear Regression to KNN: A Mathematical Deep Dive

# What is Supervised Learning?

## The Core Concept

Supervised learning is the machine learning task of learning a function that maps an input to an output based on example input-output pairs.

- **Input (X):** Features or variables (e.g., email content, house size).
- **Output (Y):** Target or label (e.g., "Spam", Price).
- **Goal:** Approximate the mapping function so well that when you have new input data (x), you can predict the output variables (Y) for that data.



# Linear Regression: The Intuition

## Fitting a Line to Data

Imagine we want to predict housing prices based on the size of the house.

Linear regression attempts to model the relationship between two variables by fitting a linear equation to observed data. One variable is considered to be an explanatory variable, and the other is considered to be a dependent variable.

Visually, we try to draw the "best fit" straight line through our scatter plot of data points.

 Scatter plot of housing prices vs size with a regression line

# Mathematical Formulation

## The Hypothesis

The function we are trying to learn is called the Hypothesis, denoted by  $h$ .

$$h_{\theta} (x) = \theta_0 + \theta_1 x$$

This is simply the equation of a straight line.

## The Parameters

The  $\theta$  values are called parameters or weights.

- $\theta_0$ : The bias unit (y-intercept).
- $\theta_1$ : The feature weight (slope).

**Goal:** Choose  $\theta_0$  and  $\theta_1$  so that  $h_{\theta} (x)$  is close to  $y$  for our training examples.

# The Cost Function (MSE)


## Measuring the Error

We need a metric to measure how well our line fits the data.


We use the **Mean Squared Error (MSE)** cost function, denoted as  $J(\theta)$ .

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m h_\theta(x^{(i)}) - y^{(i)}_i$$

Here,  $m$  is the number of training examples. We square the error to penalize large deviations. The  $1/2$  is for mathematical convenience during differentiation.



# Gradient Descent: Intuition



## Walking Down the Hill

Imagine you are standing on top of a hill (the cost function surface) and you want to get to the bottom (minimum cost) as quickly as possible.

You look around 360 degrees and ask, "If I take a baby step, in which direction will I go down the steepest?"

You take a step in that direction, and repeat the process. Eventually, you will reach the valley floor (global minimum).

# Gradient Descent Algorithm



## The Process

Repeat until convergence:

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$$



## Learning Rate

$\alpha$  (alpha) is the Learning Rate.

It controls how big of a step we take downhill. If too small, convergence is slow. If too large, we might overshoot the minimum.



## Simultaneous Update

Crucially, we must update both  $\theta_0$  and  $\theta_1$  simultaneously at each step of the iteration.

# Deriving the Update Rules

## Derivative for Bias ( $\theta_0$ )

By taking the partial derivative of the MSE cost function with respect to  $\theta_0$ :

$$\frac{\partial}{\partial \theta_0} J(\theta) = \frac{1}{m} \sum_{i=1}^m h_\theta(x^{(i)}) - y^{(i)}$$

## Derivative for Weight ( $\theta_1$ )

Similarly, for  $\theta_1$ , applying the chain rule gives us an extra  $x$  term:

$$\frac{\partial}{\partial \theta_1} J(\theta) = \frac{1}{m} \sum_{i=1}^m h_\theta(x^{(i)}) - y^{(i)} \cdot x^{(i)}$$

# Multivariate Linear Regression


## Multiple Features

When we have more than one feature (e.g., size, number of bedrooms, age of house), we denote the input as a vector  $x$ .

The hypothesis becomes:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$$

In this space, we are fitting a hyperplane rather than a simple line.



# Vectorization & Normal Equation

## Vectorized Hypothesis

To compute efficiently, we define  $x_0 = 1$  and use linear algebra:

$$h_{\theta}(x) = \theta^T x$$

This allows us to compute predictions for the entire dataset in one matrix multiplication.

## The Normal Equation

Instead of iterating with Gradient Descent, we can solve for  $\theta$  analytically:

$$\theta = (X^T X)^{-1} X^T y$$

This gives the exact minimum but is computationally expensive for very large datasets ( $O(n^3)$ ).


# Introduction to Classification

## Discrete Outputs

In classification, the variable  $y$  is discrete (e.g., 0 or 1, Spam or Not Spam).

### Why not Linear Regression?

- Linear regression predicts continuous values, which can be  $> 1$  or  $< 0$ .
- It is highly sensitive to outliers, which can shift the decision boundary drastically and incorrectly.



# Logistic Regression: The Sigmoid


## The Sigmoid Function

We want  $0 \leq h_{\theta}(x) \leq 1$ . We map the linear output through the Sigmoid function (Logistic function):


$$g(z) = \frac{1}{1 + e^{-z}}$$

Our hypothesis becomes  $h_{\theta}(x) = g(\theta^T x)$ .

This outputs the **probability** that  $y=1$ .



# The Decision Boundary



## Making a Prediction

Even though the output is probabilistic, we can enforce a hard classification.

- Predict  $y = 1$  if  $h_{\theta}(x) \geq 0.5$
- Predict  $y = 0$  if  $h_{\theta}(x) < 0.5$

This corresponds to  $\theta^T x \geq 0$ . The line where  $\theta^T x = 0$  is the decision boundary.

# Logistic Regression Cost Function

## Why not MSE?

Using Mean Squared Error with the sigmoid function results in a **non-convex** cost function with many local minima, making gradient descent unreliable.

## Log Loss (Cross-Entropy)

We use a convex cost function derived from maximum likelihood estimation:

$$\text{Cost} = -y \log(h_{\theta}(x)) - (1-y) \log(1-h_{\theta}(x))$$

This heavily penalizes confident wrong predictions (e.g., predicting 1 when actual is 0).


# K-Nearest Neighbors (KNN)

## Instance-Based Learning

KNN is a simple, non-parametric, lazy learning algorithm. It doesn't learn a "model" or coefficients like regression.

**The Logic:** "Tell me who your neighbors are, and I'll tell you who you are."

To classify a new data point, we look at the 'K' closest training examples and take a majority vote.



# KNN: Distance & Parameters

## Euclidean Distance

How do we define "closest"? The most common metric is Euclidean distance:

$$d(p, q) = \sqrt{\sum_{i=1}^n (q_i - p_i)^2}$$

## Choosing 'K'

The choice of K is crucial:

- **Small K (e.g., 1):** Sensitive to noise and outliers (High Variance).
- **Large K:** Smoother decision boundaries, but may miss local patterns (High Bias).

# Key Takeaways



## Linear Regression

Used for predicting continuous values. Optimizes MSE using Gradient Descent.



## Logistic Regression


Used for binary classification. Outputs probabilities using the Sigmoid function.



## K-Nearest Neighbors


A simple, geometric classifier based on distance to nearby training points.

# Image Sources




[https://images.prismic.io/encord/190bb719-b32f-4da7-85c7-5e3892ff2b38\\_Supervised+Learning+Flowchart---Encord.jpg?auto=compress,format](https://images.prismic.io/encord/190bb719-b32f-4da7-85c7-5e3892ff2b38_Supervised+Learning+Flowchart---Encord.jpg?auto=compress,format)

Source: [encord.com](http://encord.com)




[https://prd-api-aggregate.statcrunch.com/api/aggregation/documents/77411ZQKDW?context=results\\_image&code=&extension=png](https://prd-api-aggregate.statcrunch.com/api/aggregation/documents/77411ZQKDW?context=results_image&code=&extension=png)

Source: [www.statcrunch.com](http://www.statcrunch.com)




<https://global.discourse-cdn.com/dlai/original/3X/1/3/133adb05a8ab320f2d069daacaa20f21bbe63a9d.jpeg>

Source: [community.deeplearning.ai](http://community.deeplearning.ai)




<https://i.sstatic.net/Rq40j.png>

Source: [stackoverflow.com](http://stackoverflow.com)



<https://i.sstatic.net/Tc3YO.png>


Source: [stackoverflow.com](http://stackoverflow.com)



<https://cdn.analyticsvidhya.com/wp-content/uploads/2020/12/image-96.png>


Source: [www.analyticsvidhya.com](http://www.analyticsvidhya.com)

# Image Sources



[https://miro.medium.com/1\\*bCCcQhMjHGal89i-7i3xFw.png](https://miro.medium.com/1*bCCcQhMjHGal89i-7i3xFw.png)

Source: [medium.com](#)



[https://scipython.com/media/old\\_blog/logistic\\_regression/decision-boundary.png](https://scipython.com/media/old_blog/logistic_regression/decision-boundary.png)

Source: [scipython.com](#)



<https://insightimi.wordpress.com/wp-content/uploads/2020/03/knn-start.png>

Source: [insightimi.wordpress.com](#)